

Cloud Test: Enterprise RAG

System Architecture

Cloud Architecture Design Submission

November 2025

1. Assumptions

- Document Volume & Characteristics
- Update Frequency & Access Patterns
- User Access Patterns
- Technology Stack Familiarity

2. High-Level Architecture

- System Overview
- Data Flow Summary

3. Ingestion and Indexing Pipeline

- Document Collection
- Document Preprocessing
- Embedding Generation
- Vector Store: **Amazon OpenSearch Serverless**

4. RAG Retrieval + Response Logic

- Query Processing Flow

5. User Interface + Application Layer

- Front-End Experience
- Backend Services

6. Security Architecture

- Authentication & Authorization
- Network Security

- Audit Logging & Monitoring
- 7. Scaling Strategy
 - Document Volume Growth
 - User Concurrency Scaling
 - Performance Under Load
 - High Availability
- 8. Cost Strategy
 - Monthly Cost Breakdown (Estimated)
 - Cost Optimization Strategies
 - Cost Monitoring & Alerts
- 9. Risks, Tradeoffs, and Alternatives
 - Design Optimizations
 - Key Risks & Mitigations
 - Alternative Architectures Considered
 - Future Improvements
 - Success Metrics
- Conclusion

1. Assumptions

Document Volume & Characteristics

- **Total Documents:** ~100,000 documents accumulated over 10+ years
- **Document Types:** 60% PDFs, 25% Word documents, 10% PowerPoint, 5% emails/other
- **Average Document Size:** 2-5 MB per document
- **Total Storage:** ~300-500 GB of raw documents
- **Growth Rate:** ~10,000 new documents per year (~800/month)

Update Frequency & Access Patterns

- **Batch Uploads:** Weekly bulk uploads of 50-200 documents
- **Real-time Indexing:** Required for newly added documents within 15 minutes
- **Query Volume:** ~2,000-5,000 queries per day (avg 3 queries/user/week)
- **Peak Usage:** Business hours (9 AM - 5 PM EST), 500 concurrent users during peaks
- **Query Distribution:** 70% simple lookups, 30% complex multi-document queries

User Access Patterns

- **Total Users:** 500 employees with authenticated access
- **Active Users:** ~300 users querying monthly, ~80 users daily
- **Document Permissions:** 70% accessible to all, 20% department-restricted, 10% confidential
- **Mobile Access:** 30% of queries from mobile devices


Technology Stack Familiarity


- **Cloud Provider:** AWS (assumed based on enterprise adoption, US data residency)
- **Development Team:** 3-5 engineers familiar with Python, React, cloud services
- **Maintenance:** Preference for managed services over custom infrastructure
- **Compliance:** SOC 2, basic data residency requirements (US-only)

2. High-Level Architecture

System Overview

The RAG system follows a dual-pipeline architecture: (1) an offline ingestion pipeline for document processing and indexing, and (2) an online query pipeline for retrieval and generation.





Data Flow Summary

Ingestion Path: Document → S3 → Text Extraction → Chunking → Embedding → Vector Store + Metadata DB

Query Path: User Query → Auth → Embedding → Vector Search
→ Re-rank → LLM Context → Generated Response + Citations →
User

3. Ingestion and Indexing Pipeline

Document Collection

Source Integration

- **Primary Sources:**

- SharePoint/OneDrive API integration for existing documents
- Email integration via Microsoft Graph API for project emails
- Web upload interface for ad-hoc documents

- **Implementation:**

- AWS Lambda functions triggered by S3 events or scheduled EventBridge rules
- Initial bulk migration: AWS DataSync or custom batch scripts
- Ongoing sync: Daily scheduled Lambda pulling from SharePoint API

File Storage

- **Raw Document Storage:** S3 Standard tier

- Bucket structure: s3://company-docs-raw/{year}/{department}/{document-id}
- Lifecycle policy: Move to S3 Glacier after 1 year (documents rarely re-processed)
- Versioning enabled for document updates

Document Preprocessing

Text Extraction Pipeline

```
Lambda Function (Python 3.11, 10GB RAM, 15min timeout)
├─ PDF Processing: PyPDF2 + Tesseract OCR for scanned documents
├─ Word/PPT: python-docx, python-pptx libraries
├─ Email: email.parser library for .eml files
└─ Output: Extracted plain text + metadata (author, date, department, access level)
```

Chunking Strategy

- **Chunk Size:** 512 tokens (~400 words) with 50-token overlap
- **Method:** Semantic chunking using LangChain's RecursiveCharacterTextSplitter
 - Split on paragraphs first, then sentences
 - Preserve section headers with chunks for context
- **Metadata Preservation:** Each chunk tagged with:
 - document_id, chunk_index, source_file, page_number
 - department, created_date, author, access_level
 - document_type, project_id (if applicable)

OCR Handling

- **Scanned PDFs:** Detected via PDF metadata or text extraction failure
- **OCR Engine:** Amazon Textract for production quality
 - Cost optimization: Use async API for batch processing
 - Fallback: Tesseract OCR for cost-sensitive scenarios

Embedding Generation

Model Selection: Amazon Titan Embeddings G1 or E5-large-v2

- **Reasoning:**
 - Titan: Fully managed, AWS-native, good cost/performance
 - E5-large: Higher quality, can self-host on SageMaker for cost control
- **Dimensions:** 1024 (Titan) or 768 (E5)
- **Batch Processing:** 100 chunks per API call via SageMaker batch transform

Implementation

```
# Pseudo-code for embedding pipeline
def process_document(s3_key):
    text = extract_text(s3_key)
    chunks = chunk_document(text, size=512, overlap=50)

    embeddings = sagemaker_batch_embed(chunks)

    for chunk, embedding in zip(chunks, embeddings):
        metadata = extract_metadata(s3_key, chunk)
        store_vector(embedding, chunk.text, metadata)
        store_metadata_db(chunk.id, metadata)
```

Vector Store: Amazon OpenSearch Serverless

Configuration

- **Index Settings:**
 - Engine: k-NN with HNSW algorithm

- Similarity: Cosine similarity
- Dimensions: 1024
- M (neighbors): 16, ef_construction: 512
- **Sharding:** Auto-scaled based on data volume
- **Replicas:** 2 for high availability

Alternative: Pinecone (if cost-effective at scale)

- Serverless tier for auto-scaling
- Built-in filtering for metadata queries
- Simpler management, potential cost savings at 100k+ vectors

4. RAG Retrieval + Response Logic

Query Processing Flow

1. Query Embedding

- Same embedding model as documents (Titan/E5) for consistency
- Real-time inference via SageMaker endpoint (1-2 instances, auto-scaling)
- Latency target: <100ms for embedding generation

2. Vector Search Strategy

Initial Retrieval: k=20 documents - Cosine similarity threshold: >0.7 (configurable) - Metadata pre-filtering for access control:

```
python filter_query = { "bool": { "must": [{"term": {"access_level": user.allowed_levels}}], "should": [{"term": {"department": user.department}}] }}
```

Hybrid Search (Optional Enhancement): - Combine vector search (semantic) + BM25 (keyword) with 0.7:0.3 weighting - Improves recall for specific terms (contract IDs, project names)

3. Re-ranking Strategy

Model: Cohere Rerank API or custom cross-encoder - Input: Query + top 20 chunks - Output: Rescored and sorted top 5-10 chunks - Latency: ~200-500ms (acceptable for quality improvement)

Criteria: - Semantic relevance to query - Recency (boost documents from recent years) - Document authority (boost official templates, final reports)

4. Context Assembly

Selected Chunks: Top 3-5 after re-ranking - **Context Window:**

~3000 tokens for LLM (Claude 3 Sonnet: 200k context) -

Template: ``` Context Documents: [Document 1: {title} - {date}] {chunk_text} Source: {file_name}, Page {page_number}

[Document 2: ...]

User Question: {query}

Instructions: Answer based solely on the provided context.

Always cite sources by document title and page number. ```

5. LLM Generation

Model: Amazon Bedrock - Claude 3.5 Sonnet - Reasoning:

- Best-in-class reasoning for complex queries - Strong citation accuracy - 200k context window handles large document sets - AWS-native, simplified deployment

Inference Configuration: - Temperature: 0.1 (low for factual accuracy) - Max tokens: 1024 - Stop sequences: None (allow full responses)

Prompt Engineering:

You are a knowledge assistant for {Company}. Answer the question using ONLY the provided documents. Always cite your sources.

If the answer is not in the documents, say "I don't have information"

about that in the available documents."

Format citations as: [Document Title, Page X]

6. Citation Attachment

Post-processing: - Parse LLM response for citation markers - Validate citations against retrieved chunks - Enrich with clickable links to original documents in S3 - Format: [Healthcare Project Summary 2022, p.5] (<https://link-to-doc>)

7. Response Caching

Cache Layer: Amazon ElastiCache (Redis) - Key: Hash of (query + user_access_level) - TTL: 1 hour for frequently asked questions - Cache hit rate target: 20-30% (saves LLM costs)

5. User Interface + Application Layer

Front-End Experience

Web Application (Primary)

Technology Stack: - **Framework:** React 18 with TypeScript - **UI Library:** Material-UI or Tailwind CSS - **State Management:** React Query for API calls, Zustand for local state - **Hosting:** AWS Amplify or S3 + CloudFront

Features: - Clean chat interface (ChatGPT-like UX) - Real-time typing indicators - Inline citation cards with document previews - Search history and saved queries - Dark/light mode

Key Interactions:

The screenshot shows a user interface for an AI assistant. At the top, there is a search bar containing the text "Search: \"Healthcare projects 2022\"". Below the search bar, there is a list of items, each enclosed in a vertical box with a thin border:

- Assistant:
In 2022, we completed 8 healthcare projects focusing on EHR integrations and telemedicine platforms.
- Key challenges included:
 - HIPAA compliance requirements
 - Legacy system integrations
- Sources:
 - [1] Healthcare Summary 2022.pdf, p.3
 - [2] Project Retrospective Q4.docx

Mobile Application (Progressive Web App)

Technology: Same React app with responsive design - PWA manifest for “Add to Home Screen” - Offline support for recent queries (service workers) - Touch-optimized interface

Document Preview

- Embedded PDF viewer (PDF.js) highlighting relevant sections
- Deep links to specific pages
- Download options (with audit logging)

Backend Services

API Gateway

AWS API Gateway (REST API): - Rate limiting: 100 requests/minute per user - Request validation and sanitization - CORS configuration for web/mobile clients - Custom domain: <https://api.knowledge.company.com>

Endpoints: - POST /query - Main RAG query endpoint - GET /documents/{id} - Document retrieval - GET /history - User query history - POST /feedback - Response quality feedback

Application Orchestrator (Lambda/ECS)

Option 1: AWS Lambda (Recommended for cost) - Language: Python 3.11 - Memory: 3GB - Timeout: 30 seconds - Concurrency: 500 reserved (matches user count)

Option 2: ECS Fargate (For complex workflows) - Container:
Python FastAPI application - Auto-scaling: 2-10 tasks based on
CPU/memory - Use case: If adding complex multi-turn
conversations

Orchestration Logic:

```
async def handle_query(query, user_id):
    # 1. Authenticate and get user permissions
    user = get_user_context(user_id)

    # 2. Check cache
    cached = check_cache(query, user.access_level)
    if cached:
        return cached

    # 3. Embed query
    query_vector = await embed_query(query)

    # 4. Vector search with access filters
    candidates = search_vectors(query_vector, user.filters)

    # 5. Re-rank
    top_chunks = rerank(query, candidates)

    # 6. Generate response
    response = await llm_generate(query, top_chunks)

    # 7. Log and cache
    log_query(user_id, query, response)
    cache_response(query, response)

return response
```

Asynchronous Processing


- **SQS Queue:** For document ingestion jobs
- **Step Functions:** Orchestrate multi-stage document processing

- **EventBridge**: Schedule daily sync jobs

6. Security Architecture

Authentication & Authorization

Employee Authentication: AWS Cognito



Configuration: - **Integration:** SAML 2.0 with Azure AD/Okta (company IdP) - **MFA:** Required for all users - **Token Expiration:** 1 hour access tokens, 30-day refresh tokens - **Groups:** Synced from company directory (Engineering, Sales, HR, etc.)

Document-Level Authorization

Access Control Model:

```
Document Metadata:  
{  
    "document_id": "doc_12345",  
    "access_level": "department", # public | department | confidential
```

```

    "allowed_departments": ["Engineering", "Product"],
    "allowed_users": ["user_789"], # For confidential docs
    "created_by": "user_456"
}

User Context:
{
    "user_id": "user_123",
    "department": "Engineering",
    "groups": ["Engineering", "Managers"],
    "clearance_level": "standard" # standard | sensitive
}

```

Filter Application: - Vector search includes metadata filters (OpenSearch query filters) - Double-check: Post-retrieval validation before passing to LLM - Audit: Log when access is denied

Encryption Strategy


Data at Rest: - **S3:** SSE-S3 (AES-256) for raw documents - **OpenSearch:** Encryption at rest enabled (AWS-managed keys) - **RDS/DynamoDB:** Default encryption with KMS - **KMS Key:** Customer-managed CMK for sensitive data rotation

Data in Transit: - **TLS 1.3:** All API communications - **VPC PrivateLink:** Internal AWS service communication - **Certificate Management:** AWS Certificate Manager (ACM)

Secrets Management: - **AWS Secrets Manager:** API keys, database credentials - **IAM Roles:** Preferred over long-lived credentials - **Rotation:** Automatic 90-day rotation for DB passwords

Network Security

Network Architecture



Segmentation: - Application layer cannot directly access data layer (routed through private subnets) - No public IPs for backend services - VPC endpoints for AWS services (no internet egress)

WAF Rules (AWS WAF)

- **Rate Limiting:** 2000 requests per 5 minutes per IP
- **Geo-Blocking:** US-only access (data residency requirement)
- **OWASP Top 10:** Managed rule groups for SQLi, XSS, etc.
- **IP Reputation:** AWS managed IP reputation lists

Audit Logging & Monitoring

Audit Trail

CloudWatch Logs + S3 Archive: - **Query Logs:** User ID, timestamp, query text, results returned, access denials - **Document Access:** Who accessed which documents, when - **Authentication Events:** Login attempts, MFA challenges, session expirations - **Data Changes:** Document uploads, deletions, permission changes

Retention: - CloudWatch: 90 days (hot access for investigations) - S3 Archive: 7 years (compliance requirement)

Security Monitoring

AWS GuardDuty: Threat detection for unusual API calls or access patterns

AWS Config: Configuration compliance tracking - Rule: Ensure all S3 buckets have encryption enabled - Rule: Ensure all Lambda functions are in VPC (if required)

CloudWatch Alarms: - Spike in 403 errors (potential unauthorized access attempts) - Unusual query patterns (potential data exfiltration) - Failed authentication rate >10/minute

Compliance

Automated Evidence Collection: - **AWS Audit Manager:** Pre-configured SOC 2 frameworks - **Data Residency:** Tag all resources with region, monitor with Config - **Access Reviews:** Quarterly Cognito user group audits

7. Scaling Strategy

Document Volume Growth

Current State (Year 1)

- **Documents:** 100,000
- **Vectors:** ~500,000 (5 chunks/document avg)
- **Storage:** 500 GB raw + 50 GB vectors

Projected Growth (Year 5)

- **Documents:** 150,000 (50% growth)
- **Vectors:** 750,000
- **Storage:** 750 GB raw + 75 GB vectors

Scaling Approach

Vector Store Scaling: - **OpenSearch Serverless:** Auto-scales data nodes based on storage/query load - **Pinecone:** Automatically handles scaling to millions of vectors - **Sharding Strategy:** If self-managed, shard by department (natural partitioning)

Embedding Pipeline: - **SageMaker Endpoint:** Auto-scaling from 1-5 instances based on CloudWatch metrics - **Batch Processing:** Use SageMaker Batch Transform for bulk re-indexing (cost-effective)

Storage: - **S3:** Unlimited scaling, no action needed - **Metadata DB:** - Start: RDS PostgreSQL (db.t3.medium, 100 GB) - Scale: Upgrade to db.r5.large or migrate to DynamoDB (serverless)

User Concurrency Scaling

Current Requirement

- **Peak Users:** 500 concurrent
- **Queries:** ~5000/day (~350/hour during peaks)

Scaling Components

API Gateway: - Default limit: 10,000 requests/second (more than sufficient) - No scaling action needed

Lambda Functions: - **Reserved Concurrency:** 500 (matches peak users) - **Provisioned Concurrency:** 50 (for cold start elimination) - **Cost:** Provisioned concurrency adds ~\$500/month, reduces latency from 2s → 200ms

SageMaker Embedding Endpoint: - Auto-scaling policy:

```
yaml   ScalingPolicy:      TargetValue: 70%
      (InvocationsPerInstance)    MinCapacity: 2 instances
      MaxCapacity: 10 instances    ScaleUpCooldown: 60 seconds
      ScaleDownCooldown: 300 seconds
```

Bedrock LLM: - Fully managed, no scaling configuration - Pay-per-token pricing absorbs spikes automatically

ElastiCache (Redis): - Start: cache.t3.medium (1 node) - Scale: Add read replicas (2-3 nodes) for read-heavy workloads - Cluster mode: Enable if cache size exceeds 100 GB

Performance Under Load

Latency Targets

Component	Target	Scaling Strategy
Query Embedding	<100ms	SageMaker auto-scaling
Vector Search	<200ms	OpenSearch compute scaling
Re-ranking	<300ms	Async processing if >500ms
LLM Generation	<3s	Use Bedrock on-demand (no scaling needed)
Total E2E	<5s	Cache frequent queries for <500ms

Load Testing Plan

Tools: Locust or Artillery for API load testing

Test Scenarios: 1. **Baseline:** 100 concurrent users, 10 queries/min each 2. **Peak:** 500 concurrent users, 20 queries/min each 3. **Spike:** 1000 users for 5 minutes (2x peak)

Success Criteria: - 95th percentile latency <5 seconds - Error rate <0.1% - Cache hit rate >25%

Database Scaling

Metadata DB (RDS PostgreSQL): - **Read Replicas:** 1-2 replicas for query history, document metadata reads -

Connection Pooling: RDS Proxy (500 max connections) -

Upgrade Path: db.t3.medium → db.r5.large → Aurora Serverless v2

Alternative: DynamoDB: - Better for unpredictable scaling - Cost-effective with on-demand pricing - Trade-off: Less flexible for complex queries

High Availability

Multi-AZ Deployment

- **RDS**: Multi-AZ for automatic failover
- **OpenSearch**: 2 replicas across 3 AZs
- **ElastiCache**: Multi-AZ with automatic failover
- **Lambda**: Inherently multi-AZ

Disaster Recovery

- **RTO (Recovery Time Objective)**: 1 hour
- **RPO (Recovery Point Objective)**: 15 minutes
- **Backup Strategy**:
 - S3: Versioning + cross-region replication (to US-WEST for DR)
 - RDS: Automated backups (daily snapshots, 7-day retention)
 - OpenSearch: Automated snapshots to S3 (daily)

Uptime Target: 99.5%

- **Allowed Downtime**: ~3.6 hours/month
- **Monitoring**: CloudWatch synthetic canaries (1-minute checks)
- **Alerting**: PagerDuty integration for on-call team

8. Cost Strategy

Monthly Cost Breakdown (Estimated)

Compute & Processing

Service	Configuration	Monthly Cost
Lambda (Query handling)	350,000 invocations/month, 3GB RAM, 5s avg	\$250
Lambda (Document processing)	10,000 docs/month, 10GB RAM, 5min avg	\$200
SageMaker Endpoint (Embeddings)	2x ml.m5.large instances, 50% utilization	\$500
Bedrock (Claude 3.5 Sonnet)	5M input tokens, 1M output tokens/month	\$900
Subtotal		\$1,850

Storage

Service	Configuration	Monthly Cost
S3 (Raw documents)	500 GB Standard, 10 GB uploads/month	\$15
S3 (Processed data)	100 GB	\$3
OpenSearch Serverless	50 GB indexed data, ~500K vectors	\$800
RDS PostgreSQL	db.t3.medium, 100 GB storage	\$120
Subtotal		\$938

Networking & Caching

Service	Configuration	Monthly Cost
CloudFront	100 GB data transfer, 1M requests	\$20
API Gateway	350,000 requests/month	\$1.25
ElastiCache (Redis)	cache.t3.medium, 1 node	\$60
VPC Endpoints	3 endpoints (S3, DynamoDB, Secrets Manager)	\$22
Data Transfer	50 GB out to internet	\$5
Subtotal		\$108

Security & Monitoring

Service	Configuration	Monthly Cost
Cognito	500 MAU (monthly active users)	\$14
WAF	2 rules, 350K requests	\$10
CloudWatch Logs	20 GB ingestion, 3-month retention	\$30
GuardDuty	500 GB CloudTrail, VPC Flow Logs	\$40
Secrets Manager	5 secrets	\$2
Subtotal		\$96

Additional Services

Service	Configuration	Monthly Cost
Step Functions	10,000 state transitions/month	\$2.50
SQS	1M requests	\$0.40
Textract (OCR)	1,000 pages/month	\$15
EventBridge	10,000 events	\$0.10
Subtotal		\$18

TOTAL MONTHLY COST: ~\$3,010

Buffer for unexpected usage: ~\$500 **Target Budget:** \$8,000/month **Remaining Budget:** ~\$4,500/month (56% under budget)

Cost Optimization Strategies

Reserved Capacity (20-40% savings)

1. **RDS Reserved Instance:** 1-year commitment saves ~\$40/month
2. **SageMaker Savings Plan:** 1-year commitment saves ~\$150/month
3. **ElastiCache Reserved Nodes:** 1-year commitment saves ~\$20/month
4. **Total Savings:** ~\$210/month = \$2,520/year

Auto-Scaling & Right-Sizing

1. **SageMaker Endpoint:** Scale to 1 instance during off-hours (6 PM - 8 AM)
 - Savings: ~\$150/month
2. **Lambda Provisioned Concurrency:** Only during business hours (9 AM - 6 PM)
 - Savings: ~\$300/month (if using provisioned concurrency)
3. **RDS Instance Scheduler:** Stop dev/test instances on weekends
 - Not applicable for production, but consider for staging environments

Managed vs. Self-Hosted Tradeoffs

Option 1: Replace OpenSearch Serverless with Pinecone

Serverless - Pinecone: ~\$70/month for 500K vectors (starter tier) - Savings: ~\$730/month - Trade-off: Vendor lock-in, data egress costs

Option 2: Replace SageMaker with Self-Hosted

Embeddings (ECS) - ECS Fargate: 2x small tasks (~\$100/month) - Savings: ~\$400/month - Trade-off: Maintenance overhead, slower cold starts

Option 3: Use Bedrock Claude 3 Haiku (instead of Sonnet)

- Haiku: 80% cheaper for simpler queries - Hybrid approach: Route simple queries to Haiku, complex to Sonnet - Savings: ~\$500/month (if 60% queries can use Haiku)

Caching Strategy (Cost Reduction)

- **ElastiCache Hit Rate:** 30% → saves 30% of LLM API calls
 - LLM cost reduction: ~\$270/month
 - Cache cost: \$60/month
 - **Net Savings:** \$210/month

Data Lifecycle Management

1. S3 Lifecycle Policies:

- Move raw documents to Glacier after 1 year: saves ~\$10/month (minimal)
- Delete temporary processing files after 7 days

2. CloudWatch Logs:

- Reduce retention from 90 days to 30 days for non-audit logs: saves ~\$10/month

3. OpenSearch Index Lifecycle:

- Archive older document versions (if versioning enabled)

Alternative: Serverless-First Architecture

Replace RDS with DynamoDB: - DynamoDB On-Demand:

~\$50-100/month (vs \$120 RDS) - Savings: ~\$20-70/month -

Trade-off: NoSQL schema design, less flexible querying

Cost Monitoring & Alerts

AWS Cost Explorer: - Daily cost tracking by service - Budget alerts at 80%, 100%, 120% of \$3,500/month

Anomaly Detection: - AWS Cost Anomaly Detection for unexpected spikes - Alert on >20% day-over-day increases

Cost Allocation Tags: - Tag all resources: Project:KnowledgeRAG, Environment:Production, Department:Engineering - Track costs by department for chargeback

9. Risks, Tradeoffs, and Alternatives

Design Optimizations

Primary Focus

1. **Security & Compliance:** Document-level access control, audit logging, US data residency
2. **Cost-Effectiveness:** Under budget by 56%, managed services for lower maintenance
3. **Developer Velocity:** Familiar stack (AWS, Python, React), managed services reduce ops burden

Intentionally Deprioritized

1. Multi-Cloud Strategy

- **Why:** Adds complexity, increases costs, not required by business constraints
- **Risk:** Vendor lock-in to AWS
- **Mitigation:** Use portable abstractions (LangChain), avoid AWS-specific APIs where possible

2. Advanced RAG Techniques

- **Not Implemented:** Multi-query generation, query decomposition, agentic workflows
- **Why:** Simple RAG meets 80% of use cases, added complexity not justified yet
- **Future:** Iterate based on user feedback and failure analysis

3. Real-Time Collaboration Features

- **Not Implemented:** Multi-user chat sessions, shared workspaces
- **Why:** Not in requirements, adds architectural complexity
- **Future:** Phase 2 feature if demand exists

4. Advanced Analytics

- **Not Implemented:** Query trend analysis, document popularity tracking
- **Why:** Focus on core functionality first
- **Future:** Add once baseline is stable

Key Risks & Mitigations

Technical Risks

Risk	Impact	Probability	Mitigation
Poor retrieval quality	High	Medium	<ul style="list-style-type: none">- Extensive testing with real queries- A/B test chunking strategies- Implement user feedback loop- Add re-ranking layer
LLM hallucination	High	Medium	<ul style="list-style-type: none">- Low temperature (0.1)- Strict prompt engineering- Citation validation- User feedback: “Was this helpful?”
High latency (>10s)	Medium	Low	<ul style="list-style-type: none">- Aggressive caching- Async query processing

Risk	Impact	Probability	Mitigation
			<ul style="list-style-type: none"> - Optimize vector search parameters - Provision concurrency for Lambda
OCR accuracy for scanned docs	Medium	Medium	<ul style="list-style-type: none"> - Use Textract (high accuracy) - Flag low-confidence extractions - Manual review queue for critical docs
Access control bypass	High	Low	<ul style="list-style-type: none"> - Defense in depth: metadata filters + post-retrieval checks - Penetration testing - Regular access audits

Business Risks

Risk	Impact	Probability	Mitigation
User adoption resistance	High	Medium	<ul style="list-style-type: none"> - Extensive user training - Gradual rollout (pilot → full) - Champions program - Show value early (quick wins)
Document quality issues	Medium	High	<ul style="list-style-type: none"> - Data quality assessment during migration

Risk	Impact	Probability	Mitigation
			<ul style="list-style-type: none"> - Metadata enrichment workflows - User reporting for bad results
Cost overruns	Medium	Low	<ul style="list-style-type: none"> - Conservative estimates with 50% buffer - Cost alerts at 80% - Auto-scaling limits to prevent runaway costs
Compliance violations	High	Low	<ul style="list-style-type: none"> - Legal review of architecture - Regular compliance audits - Automated evidence collection

Operational Risks

Risk	Impact	Probability	Mitigation
On-call burden	Low	Medium	<ul style="list-style-type: none"> - Comprehensive monitoring - Runbook documentation - Auto-remediation for common issues - Managed services reduce ops load
Vendor API limits	Medium	Low	<ul style="list-style-type: none"> - Monitor Bedrock quotas - Request limit increases proactively - Implement graceful degradation

Risk	Impact	Probability	Mitigation
Data loss	High	Low	<ul style="list-style-type: none"> - Multi-AZ deployments - Cross-region S3 replication - Regular backup testing - Immutable audit logs

Alternative Architectures Considered

1. Self-Hosted LLM (e.g., Llama 3, Mistral)

Pros: - No per-token costs, predictable pricing - Full control over model behavior - Data never leaves VPC

Cons: - High upfront GPU costs (~\$2,000+/month for p3.2xlarge) - Ops burden: model updates, fine-tuning, monitoring - Lower quality vs. Claude 3.5 Sonnet - Harder to scale for bursty traffic

Verdict: Not cost-effective at this scale, revisit if query volume exceeds 50K/day

2. Hybrid Search (Vector + Traditional Search)

Pros: - Better recall for specific terms (project IDs, contract numbers) - Complements semantic search

Cons: - Adds complexity (2 indexes to maintain) - Minimal improvement for narrative queries

Verdict: Consider as Phase 2 optimization after measuring retrieval quality

3. Fine-Tuned Embedding Model

Pros: - Domain-specific embeddings for consulting terminology -
Potentially better retrieval accuracy

Cons: - Requires labeled data (1000+ query-document pairs) -
Training and hosting costs - Maintenance overhead

Verdict: Not justified for v1, revisit if retrieval quality is poor

4. Multi-Modal RAG (Images, Charts, Tables)

Pros: - Extract information from visuals in PowerPoints, PDFs -
More comprehensive knowledge coverage

Cons: - Requires vision-capable LLM (Claude 3.5 supports this) -
Text extract table extraction adds cost - Increased complexity

Verdict: Phase 2 feature, implement if users request it

5. GraphRAG (Knowledge Graphs)

Pros: - Better for relationship queries ("Who worked on X
project?") - Explainable retrieval paths

Cons: - Requires entity extraction and graph construction -
Complex to implement and maintain - Minimal benefit for
document-centric queries

Verdict: Not suitable for this use case, better for knowledge
graphs with structured data

Future Improvements

Phase 2 (Months 6-12)

1. **Multi-Turn Conversations:** Maintain chat history, context-aware follow-ups
2. **Query Suggestions:** Auto-complete based on popular queries
3. **Advanced Filters:** Filter by date range, document type, department
4. **User Feedback Loop:** “Helpful/Not Helpful” buttons → improve retrieval
5. **Analytics Dashboard:** Popular queries, low-quality responses, usage trends

Phase 3 (Year 2+)

1. **Proactive Insights:** Weekly digests of relevant new documents
2. **Document Summarization:** Auto-generate summaries for long documents
3. **Hybrid Search:** Combine vector + keyword search for specific queries
4. **Multi-Lingual Support:** If expanding to international offices
5. **Fine-Tuned Models:** Custom embeddings and LLMs for domain expertise

Success Metrics

Technical KPIs

- **Retrieval Accuracy:** >85% of queries return relevant documents (measured via user feedback)

- **Latency:** 95th percentile <5 seconds
- **Uptime:** >99.5% (within SLA)
- **Cache Hit Rate:** >25%

Business KPIs

- **User Adoption:** >60% of employees use system monthly within 6 months
- **Query Volume:** 5,000+ queries/month by month 3
- **Time Savings:** 30% reduction in time spent searching for documents (user survey)
- **Cost Efficiency:** Remain under \$4,000/month in Year 1

Quality KPIs

- **Citation Accuracy:** >95% of citations point to correct sources
 - **User Satisfaction:** >4.0/5.0 average rating
 - **False Positive Rate:** <10% of responses marked “Not Helpful”
-

Conclusion

This RAG system architecture balances security, scalability, and cost-effectiveness for a mid-size consulting firm. By leveraging AWS managed services (Bedrock, OpenSearch Serverless, Cognito), we minimize operational overhead while maintaining enterprise-grade security controls.

The design prioritizes: 1. **Document-level access control** to protect confidential information 2. **Cost efficiency** (56% under budget with room to scale) 3. **Developer productivity** (familiar stack, managed services) 4. **Iterative improvement** (phased rollout, feedback loops)

With comprehensive monitoring, audit trails, and a clear scaling path, the system is production-ready and positioned for long-term success.